


PyData NYC 2024

Udisha Dutta Chowdhury

Joint work with Abhishek Murthy,
Schneider Electric



***Adopting Open-Source Tools for Time Series
Forecasting: Opportunities and Pitfalls***

About Me

Udisha Dutta Chowdhury

- Master's Student in **Computer System's Engineering** at **Northeastern University**, Boston, MA.
- **Teaching Assistant** for Machine Learning for IoT Systems Course at **Northeastern University**.
- **Data Science Intern** at **Schneider Electric**, Andover, MA.
- Cyber Security Intern and Solution Delivery Analyst at **Deloitte USI**, Bangalore, India.
- Undergrad in **Electronics and Communication Engineering**, **PES University**, Bangalore, India.



Outline of the Talk



Introduction to Time Series Forecasting



Open-Source Tools



3 critical aspects of forecasting libraries

Data Understanding
Data Preparation
Backtesting



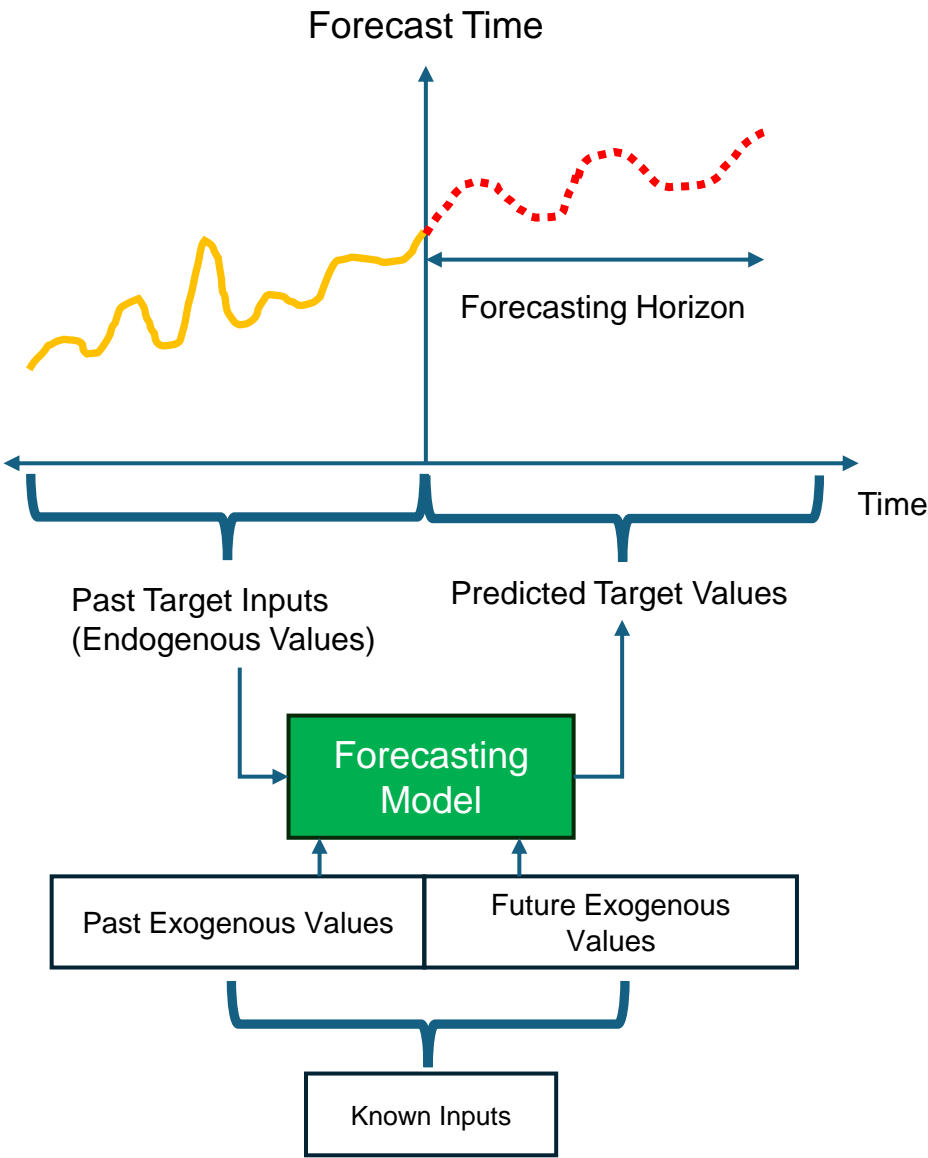
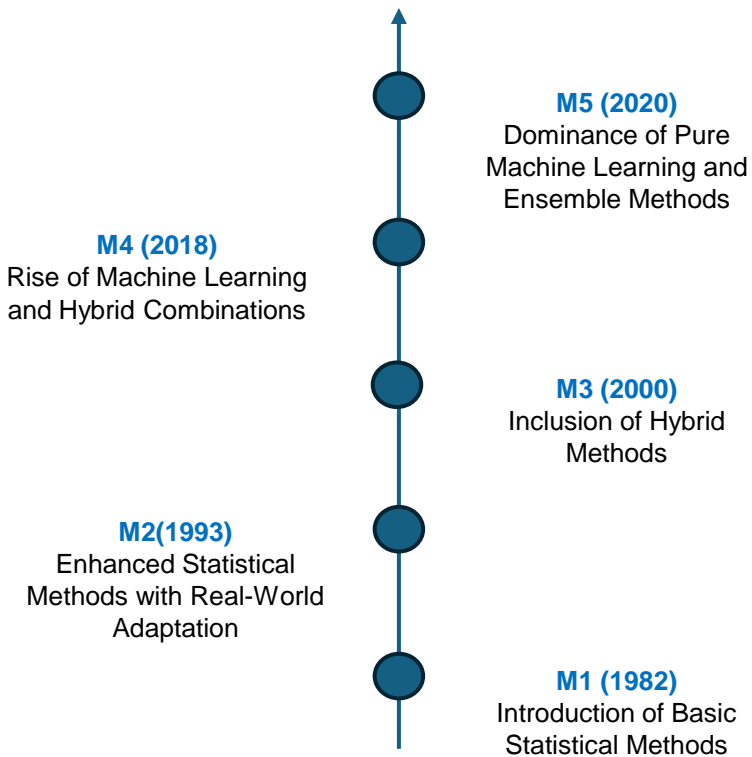
Conclusion and Key Takeaways

Time Series Forecasting

Evolution of Forecasting Models


M Competitions benchmark forecasting models, showcasing advancements from traditional statistical methods to modern machine learning approaches

Milestones in the M Competitions



Common Stages in a Forecasting Project

- **Business Understanding:**
 - Define the forecasting objective: What do we need to predict, and why is it important?
- **Data Understanding:**
 - Explore time-series data: Identify trends, seasonality, and patterns.
- **Data Preparation:**
 - Preprocess data for forecasting: Handle missing values, NaNs, duplicate values.
- **Modeling:**
 - Apply forecasting methods: Choose models (e.g., ARIMA, neural networks) based on data characteristics.
- **Evaluation (Backtesting):**
 - Test model accuracy under realistic settings.
- **Deployment:**
 - Rollout the model in production.



Open-Source Tools for Time Series Forecasting

Advantages of Open-Source Solutions

- Accessibility and Cost-Effectiveness
- Flexibility and Customizability
- Community-Driven Innovation
- Integration with Existing Ecosystems



Neural Prophet

Growing Popularity of SKtime and Skforecast

☆ 7.9k stars
👁 106 watching
🔗 1.4k forks

[Report repository](#)

Releases 81

📦 v0.34.0 Latest
last week

[+ 80 releases](#)

Sponsor this project

 sktime sktime 

 opencollective.com/sktime

[Learn more about GitHub Sponsors](#)

Used by 3.3k

 + 3,315

☆ 1.1k stars
👁 10 watching
🔗 131 forks



[Report repository](#)



Releases 28

📦 v0.13.0 Latest
on Aug 1

[+ 27 releases](#)

Sponsor this project

 JoaquinAmatRodrigo Joaquín A... 

 JavierEscobarOrtiz Javier Escoba... 

<https://www.buymeacoffee.com/skfore...>

<https://www.paypal.com/donate/?hoste...>

[Learn more about GitHub Sponsors](#)

Packages

No packages published

Used by 328

 + 320

- GitHub Stars and Repository Activity
- Community Engagement and Usage Trends



3- Critical Aspects for Forecasting Libraries



Data Understanding: How well does the library support Exploratory Data Analysis (EDA)?



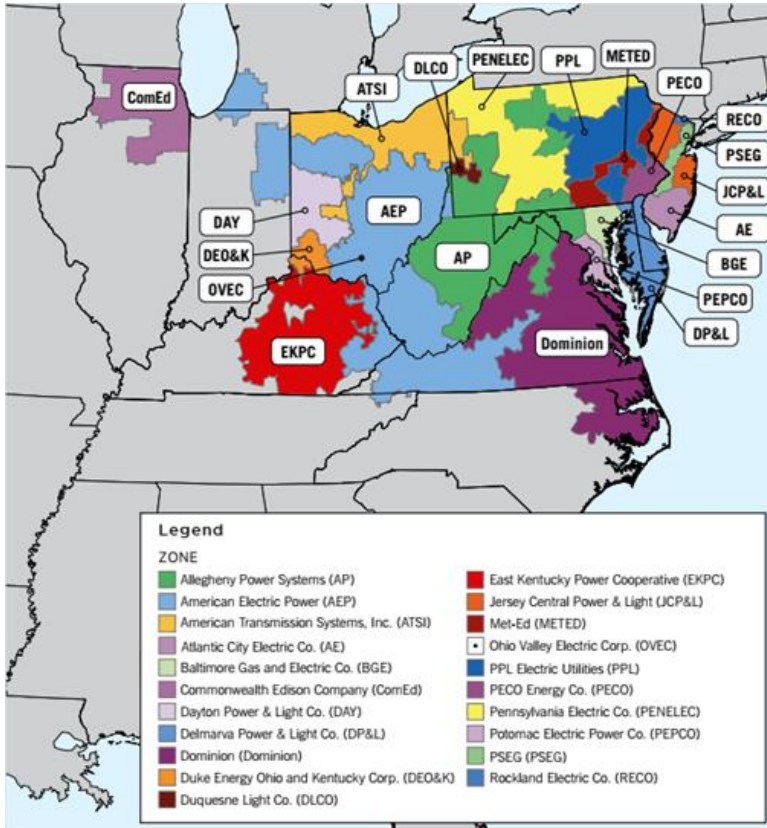
Data Preparation: How robust and intuitive are the tool's preprocessing capabilities for handling quality issues, like missing values, NaNs, duplicate data, and exogenous variables?



Backtesting: How effective and scalable are the library's modeling and evaluation capabilities for forecasting algorithms?

Dataset used for the experiments

PJM Hourly Energy Consumption Data



- Independent nonprofit organization that manages the electric transmission system for a large region
- Hourly power consumption data comes from PJM's website and are in megawatts (MW).
- Data used : 125 days of hourly power consumption of Duquesne Light Company (DUQ)


Data Understanding : Forecasting Libraries accelerate EDA



Viz. trends with train-test split

`plot_series()`


```
plot_series(y_train, y_test, labels=["Training", "Testing"], markers=['.', '.'], ax=ax)
```



Plot one or more lagged versions of the time-series data


`plot_lags()`

Plot of series against lags 6



`plot_lags(y_train, lags=6)`

Plot of series against lags 24



`plot_lags(y_train, lags=24)`

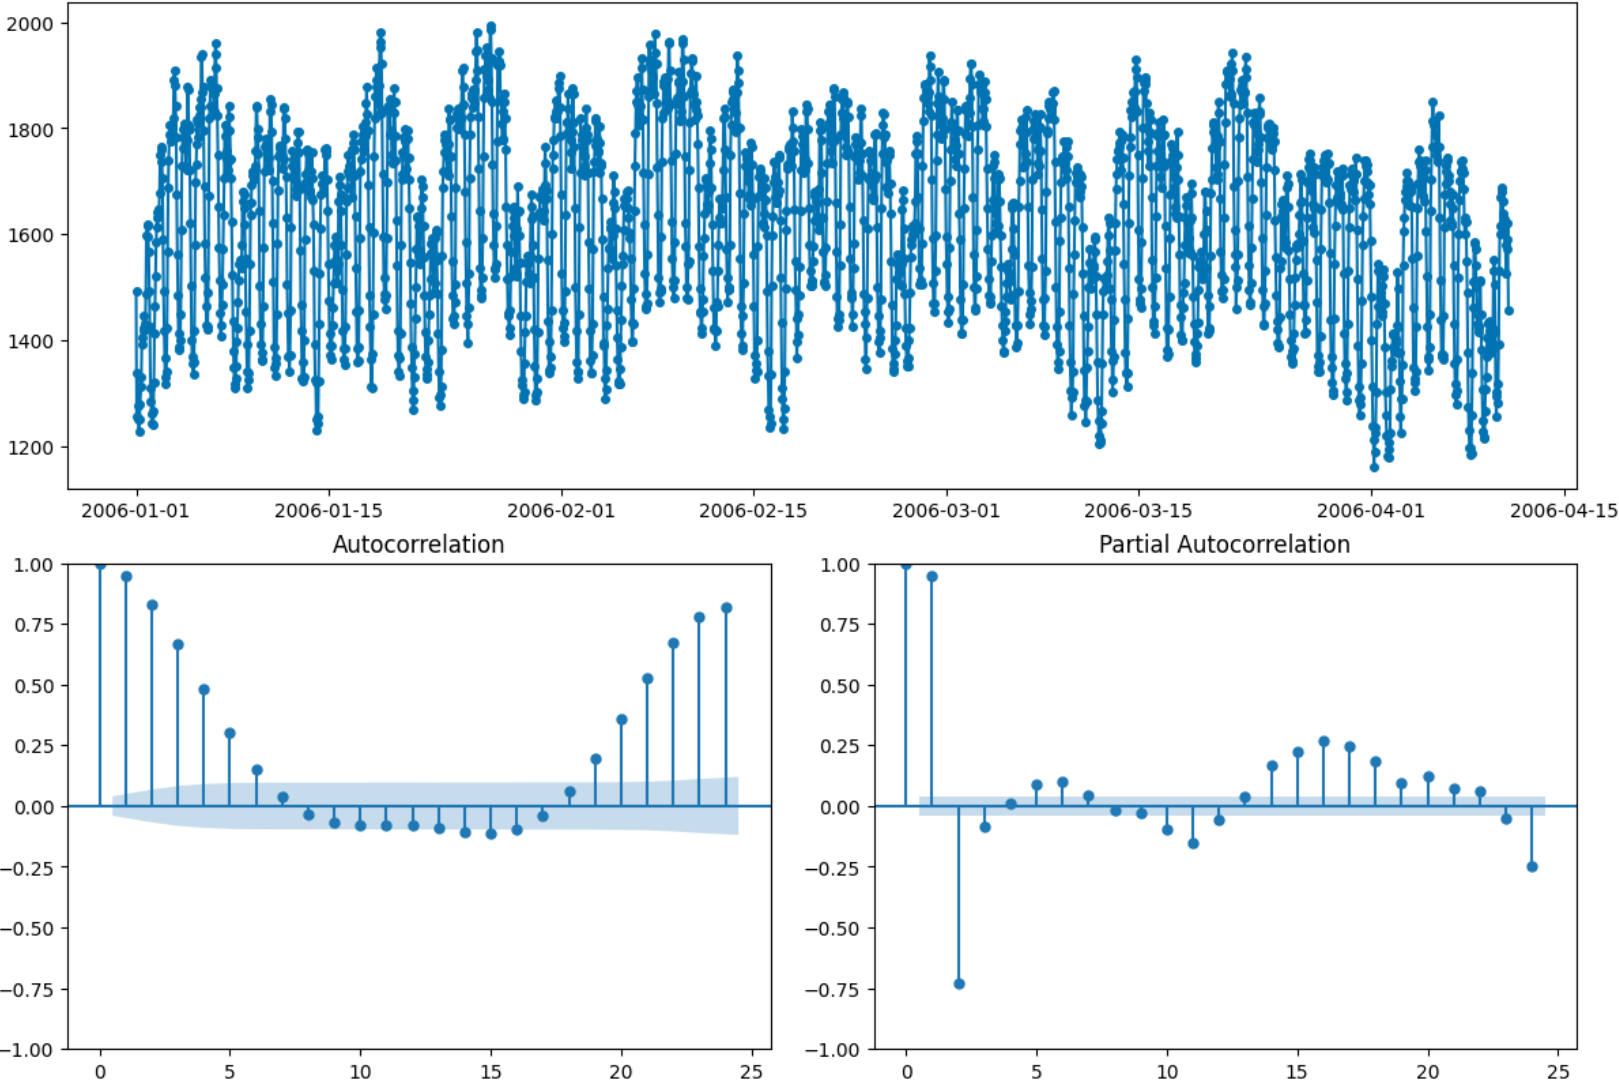
Data Understanding



Correlation (auto & partial)

`plot_correlations()`

`plot_correlations(y_train)`



Data Preparation: Forecasting Libraries offer pre-processing features

time	y	time	y	time	y	time	y
t_0	y_0	t_0	y_0	t_0	y_0	t_0	y_0
$t_0 + \delta$	y_1	$t_0 + \delta$	y_1	$t_0 + \delta$	y_1	$t_0 + \delta$	y_1
$t_0 + 2\delta$	y_2	$t_0 + 2\delta$	NaN	$t_0 + 3\delta$	y_3	$t_0 + \delta$	y_1
$t_0 + 3\delta$	y_3	$t_0 + 3\delta$	y_3	$t_0 + 4\delta$	y_4	$t_0 + 2\delta$	y_2
.		.		.		.	
.		.		.		.	
.		.		.		.	
Nominal		Missing data		Irregular data		Duplicate data	

Data quality issues are expected in the training and testing datasets.

The forecasting library must help the data scientist analyze and mitigate the issues through:

- Missing value imputation
- Outlier detection and removal
- Detrending and seasonality adjustment
- Resampling irregular timestamps to fixed frequency
- Time-based train-test splitting utilities

Data Preparation



Missing Value Imputation

```
from sktime.transformations.series.impute import Imputer
```


```
imputer = Imputer(method="mean")  
data_imputed = imputer.fit_transform(y_train)
```

```
Original Data with Missing Values:  
2006-01-01 00:00:00    0.496714  
2006-01-01 01:00:00         NaN  
2006-01-01 02:00:00    1.006138  
2006-01-01 03:00:00         NaN  
2006-01-01 04:00:00    2.295015  
Freq: H, Name: Value, dtype: float64  
  
Imputed Data:  
2006-01-01 00:00:00    0.496714  
2006-01-01 01:00:00   50.037261  
2006-01-01 02:00:00    1.006138  
2006-01-01 03:00:00   50.037261  
2006-01-01 04:00:00    2.295015  
Freq: H, Name: Value, dtype: float64
```

Outlier Detection and Removal

```
from sktime.transformations.series.outlier_detection  
import HampelFilter
```

```
hampel_filter = HampelFilter(window_length=24,  
n_sigma=3)  
outliers = hampel_filter.fit_transform(y_train)
```




Data Preparation




Detrending and Seasonality Adjustment

```
from sktime.transformations.series.detrend import  
Detrender  
  
detrender = Detrender()  
y_train_detrended = detrender.fit_transform(y_train)
```



Time-based train-test split

```
from sktime.split import temporal_train_test_split  
  
data = df['2006-01-01':'2006-12-31'][::(24*125)]  
y_train, y_test =  
    temporal_train_test_split(data, test_size=0.2)
```



Data Preparation



Skforecast works with pandas for data preprocessing

Handle Outliers using Rolling Statistics

```
roll_mean = data_clean.rolling(window=5, center=True).mean()  
roll_std = data_clean.rolling(window=5, center=True).std()
```

Detrending Using Differentiation

```
ts_diff = ts_no_outliers.diff().dropna()
```

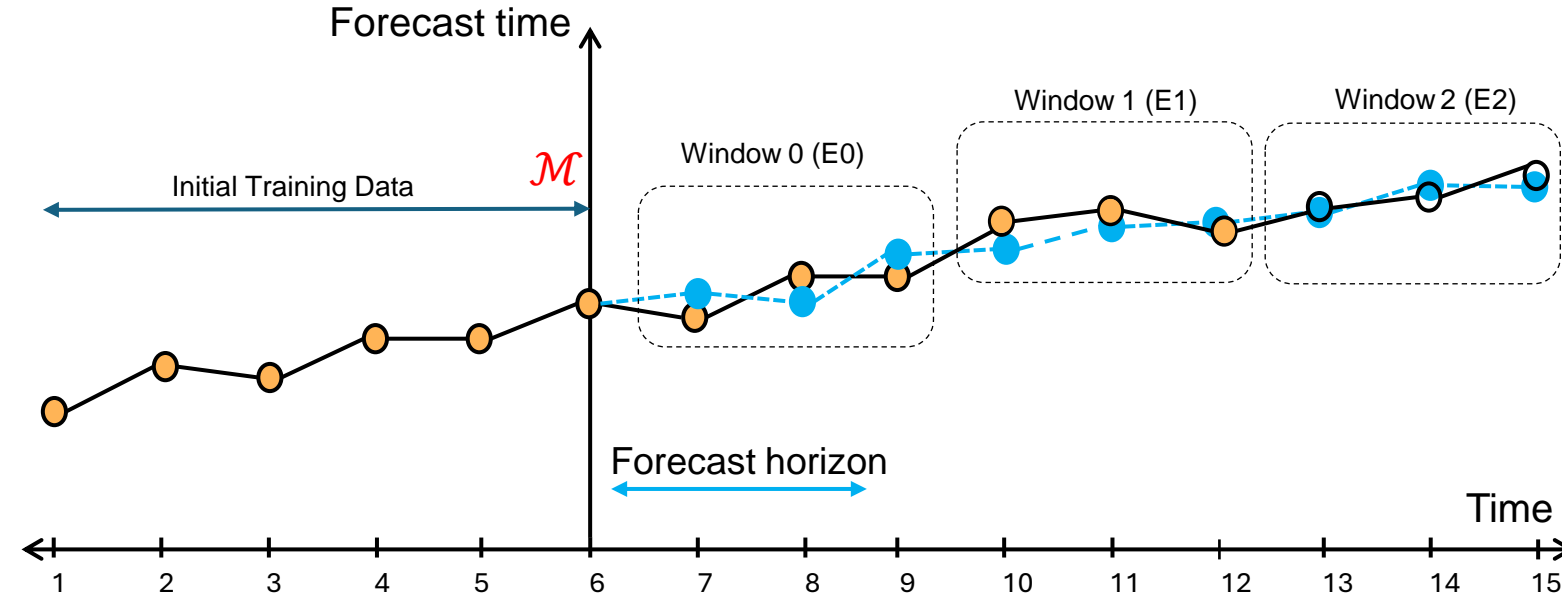
Resampling Irregular Timestamps

```
data_resampled = data.resample('H').mean()
```

Handle missing values

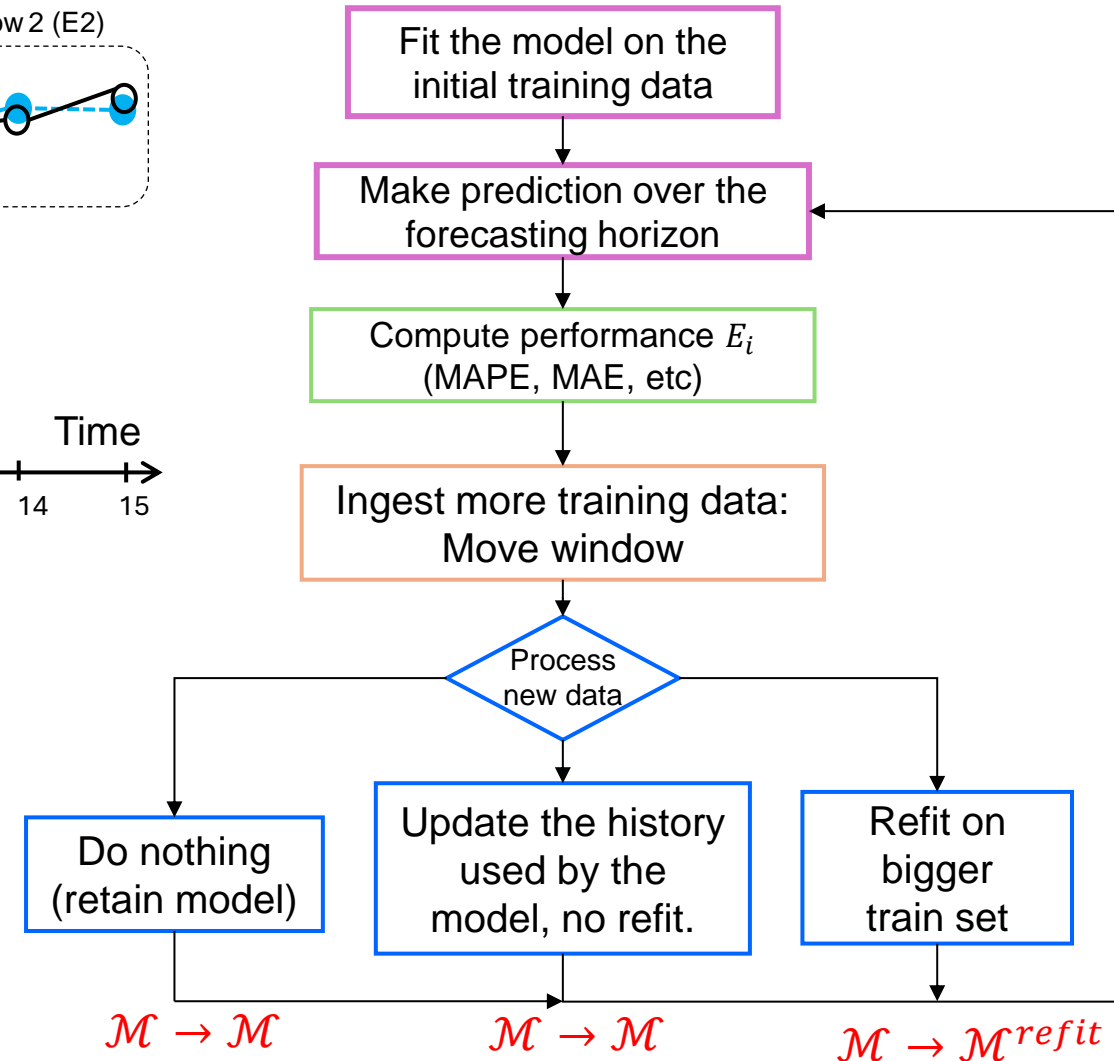
```
data_interpolated = data.interpolate(method='linear')
```


Backtesting: Forecasting libraries help scale experimentation



Forecasting libraries offer

- wide variety of modeling choices by interfacing/ implementing popular algorithms (sci-kit, XGBoost, etc)
- Standardized implementation of forecasting metrics
- Windowing functions to replicate forecasting scenarios: sliding, expanding, etc
- Model update/refit capabilities to replicate forecasting scenarios.



Backtesting: Setting Up Your Experiment Parameters



Modeling Choices

```
regressor =  
    RandomForestRegressor(n_estimators=250,  
                          max_depth=10, random_state=123)  
  
forecaster = make_reduction(regressor,  
                            window_length=6, strategy='recursive')
```



```
regressor =  
    RandomForestRegressor(n_estimators=250,  
                          max_depth=10, random_state=123)  
  
forecaster = ForecasterAutoreg(  
    regressor = regressor,  
    lags = 6)
```

Sliding/Expanding Window Features

```
cv = ExpandingWindowSplitter(initial_window=240,  
                             step_length=24,  
                             fh=np.arange(1, 25))
```

```
results = evaluate(forecaster=forecaster,  
                  y=y_train,  
                  cv=cv,  
                  return_data=True,  
                  strategy='refit',  
                  scoring=mean_absolute_percentage_error)
```

Refit/Update Features

Performance Metrics

```
metric, predictions_backtest= backtesting_forecaster(  
    forecaster = forecaster,  
    y           = y_train["DUQ_MW"],  
    initial_train_size = 240,  
    fixed_train_size  = False,  
    steps              = 24,  
    metric= mean_absolute_percentage_error,  
    refit      = True,  
    verbose    = True,  
    show_progress = True)
```

Backtesting Experiment



Training set, Testing set and Predictions for each window

test_MeanAbsolutePercentageError	fit_time	pred_time	len_train_window	cutoff	y_train	y_test	y_pred
0	0.030418	0.527055	0.110967	240	2006-01-10 23:00:00	DUQ_MW 2006-01-01 00:00:0...	DUQ_MW 2006-01-11 00:00:0...
1	0.053364	0.620258	0.102912	264	2006-01-11 23:00:00	DUQ_MW 2006-01-01 00:00:0...	DUQ_MW 2006-01-12 00:00:0...
2	0.038886	0.535657	0.099865	288	2006-01-12 23:00:00	DUQ_MW 2006-01-01 00:00:0...	DUQ_MW 2006-01-13 00:00:0...
3	0.067248	0.590521	0.106738	312	2006-01-13 23:00:00	DUQ_MW 2006-01-01 00:00:0...	DUQ_MW 2006-01-14 00:00:0...
4	0.063785	0.677798	0.099203	336	2006-01-14 23:00:00	DUQ_MW 2006-01-01 00:00:0...	DUQ_MW 2006-01-15 00:00:0...
...
85	0.027762	3.444988	0.113939	2280	2006-04-05 23:00:00	DUQ_MW 2006-01-01 00:00:0...	DUQ_MW 2006-04-06 00:00:0...
86	0.072161	3.599273	0.110180	2304	2006-04-06 23:00:00	DUQ_MW 2006-01-01 00:00:0...	DUQ_MW 2006-04-07 00:00:0...
87	0.207004	3.467539	0.110451	2328	2006-04-07 23:00:00	DUQ_MW 2006-01-01 00:00:0...	DUQ_MW 2006-04-08 00:00:0...
88	0.204000	3.575651	0.104502	2352	2006-04-08 23:00:00	DUQ_MW 2006-01-01 00:00:0...	DUQ_MW 2006-04-09 00:00:0...
89	0.086414	3.566354	0.097486	2376	2006-04-09 23:00:00	DUQ_MW 2006-01-01 00:00:0...	DUQ_MW 2006-04-10 00:00:0...

90 rows × 8 columns

End of each training window

Length of each window

Prediction time for each window

Training Time for each window

Error for each window

Backtesting Experiment



skforecast

`backtesting_forecaster()` -> returns all the predicted values and the overall performance metric value

Information of backtesting process

```
-----
Number of observations used for initial training: 240
Number of observations used for backtesting: 2160
  Number of folds: 90
  Number skipped folds: 0
  Number of steps per fold: 24
  Number of steps to exclude from the end of each train set before test (gap): 0
```

Fold: 0

```
  Training: 2006-01-01 00:00:00 -- 2006-01-10 23:00:00 (n=240)
  Validation: 2006-01-11 00:00:00 -- 2006-01-11 23:00:00 (n=24)
```

Fold: 1

```
  Training: 2006-01-01 00:00:00 -- 2006-01-11 23:00:00 (n=264)
  Validation: 2006-01-12 00:00:00 -- 2006-01-12 23:00:00 (n=24)
```

Fold: 2

```
  Training: 2006-01-01 00:00:00 -- 2006-01-12 23:00:00 (n=288)
  Validation: 2006-01-13 00:00:00 -- 2006-01-13 23:00:00 (n=24)
```

Fold: 3

```
  Training: 2006-01-01 00:00:00 -- 2006-01-13 23:00:00 (n=312)
  Validation: 2006-01-14 00:00:00 -- 2006-01-14 23:00:00 (n=24)
```

Fold: 4

```
  Training: 2006-01-01 00:00:00 -- 2006-01-14 23:00:00 (n=336)
  Validation: 2006-01-15 00:00:00 -- 2006-01-15 23:00:00 (n=24)
```

Fold: 5

...

Fold: 89

```
  Training: 2006-01-01 00:00:00 -- 2006-04-09 23:00:00 (n=2376)
  Validation: 2006-04-10 00:00:00 -- 2006-04-10 23:00:00 (n=24)
```

Training set, Testing set and the size for each fold can be seen when `Verbosity` is set to `True`

predictions_backtest

✓ 0.0s

	pred
2006-01-11 00:00:00	1520.168000
2006-01-11 01:00:00	1449.010000
2006-01-11 02:00:00	1398.115600
2006-01-11 03:00:00	1380.080162
2006-01-11 04:00:00	1355.504533
...	...
2006-04-10 19:00:00	1690.002351
2006-04-10 20:00:00	1676.797505
2006-04-10 21:00:00	1671.707992
2006-04-10 22:00:00	1678.172647
2006-04-10 23:00:00	1682.843514

2160 rows × 1 columns

metric

✓ 0.0s

mean_absolute_percentage_error



0 0.062876



Backtesting Experiment

Backtesting Parameters:

window_movement = expanding
initial_data = 240
fh = 24
window_stride = 24
metric = MAPE
refit_update = refit
window_length = 24



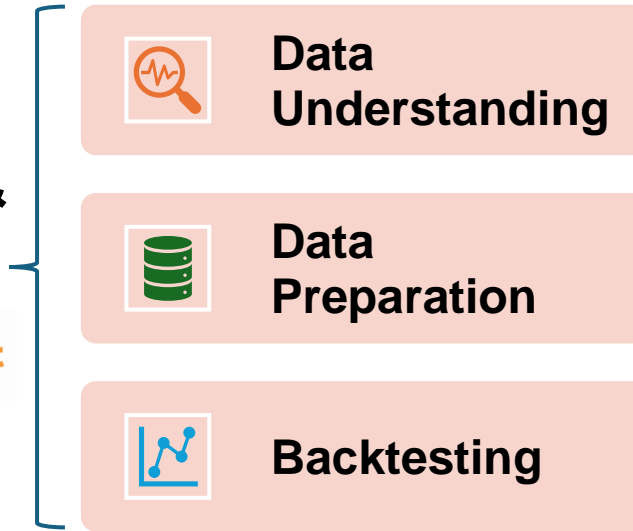


		
Experiment run time (ms)	117357.61	33883.11
Experiment Metric (MAPE)	0.088353	0.062876
Window length vs Training time		-
Window Length vs Prediction time		-



Conclusion and Key Takeaways

Minimize boilerplate code
accelerate development &
simplify maintenance



- SKTime and SKForecast: two popular forecasting libraries. SKForecast focuses on a core set of features; SKTime is all encompassing.
- **Data Understanding:** SKTime hooks into `statsmodels` to provide several visualization functions.
- **Data Preparation:** SKTime provides several utilities, whereas SKForecast expects the data scientist to use Pandas preprocessing functions.
- **Backtesting:** Despite limited functionality, SKForecast is quite fast. SKTime provides more options and fine-grained stats.